

CS 446: Machine Learning

Homework 2

Due on Tuesday, January 30, 2018, 11:59 a.m. Central Time

1. [6 points] Linear Regression Basics

Consider a linear model of the form $\hat{y}^{(i)} = \mathbf{w}^\top \mathbf{x}^{(i)} + b$, where $\mathbf{w}, \mathbf{x} \in \mathbb{R}^K$ and $b \in \mathbb{R}$. Next, we are given a training dataset, $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}$ denoting the corresponding input-target example pairs.

- (a) What is the loss function, \mathcal{L} , for training a linear regression model? (Don't forget the $\frac{1}{2}$)

Your answer: $\mathcal{L} = \frac{1}{|\mathcal{D}|} \sum_i^{|\mathcal{D}|} \frac{1}{2} (y^{(i)} - \hat{y}^{(i)})^2$

- (b) Compute $\frac{\partial \mathcal{L}}{\partial \hat{y}^{(i)}}$.

Your answer: $\frac{\partial \mathcal{L}}{\partial \hat{y}^{(i)}} = \frac{-1}{|\mathcal{D}|} (y^{(i)} - \hat{y}^{(i)})$

- (c) Compute $\frac{\partial \hat{y}^{(i)}}{\partial \mathbf{w}_k}$, where \mathbf{w}_k denotes the k^{th} element of \mathbf{w} .

Your answer: $\frac{\partial \hat{y}^{(i)}}{\partial \mathbf{w}_k} = \mathbf{x}_k^{(i)}$

- (d) Putting the previous parts together, what is $\nabla_{\mathbf{w}} \mathcal{L}$?

Your answer: $\nabla_{\mathbf{w}} \mathcal{L} = \frac{-1}{|\mathcal{D}|} \sum_i^{|\mathcal{D}|} (y^{(i)} - (w^T \mathbf{x}^{(i)} + b)) \mathbf{x}^{(i)}$

- (e) Compute $\frac{\partial \mathcal{L}}{\partial b}$.

Your answer: $\frac{\partial \mathcal{L}}{\partial b} = \frac{-1}{|\mathcal{D}|} \sum_i^{|\mathcal{D}|} (y^{(i)} - (w^T \mathbf{x}^{(i)} + b))$

- (f) For convenience, we group \mathbf{w} and b together into \mathbf{u} , then we denote $\mathbf{z} = [\mathbf{x} \ 1]$. (i.e. $\hat{y} = \mathbf{u}^\top [x, 1] = \mathbf{w}^\top x + b$). What are the optimal parameters $\mathbf{u}^* = [\mathbf{w}^*, b^*]$? Use the notation $\mathbf{Z} \in \mathbb{R}^{|\mathcal{D}| \times (K+1)}$ and $\mathbf{y} \in \mathbb{R}^{|\mathcal{D}|}$ in the answer. Where, each row of \mathbf{Z}, \mathbf{y} denotes an example input-target pair in the dataset.

Your answer: $\mathbf{u}^* = (\mathbf{Z}^T \mathbf{Z})^{-1} \mathbf{Z}^T \mathbf{y}$

2. [2 points] Linear Regression Probabilistic Interpretation

Consider that the input $x^{(i)} \in \mathbb{R}$ and target variable $y^{(i)} \in \mathbb{R}$ to have the following relationship.

$$y^{(i)} = w \cdot x^{(i)} + \epsilon^{(i)}$$

where, ϵ is independently and identically distributed according to a Gaussian distribution with zero mean and unit variance.

- (a) What is the conditional probability $p(y^{(i)} | x^{(i)}, w)$.

Your answer: $p(y^{(i)} | x^{(i)}, w) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(y^{(i)} - w x^{(i)})^2}{2}\right)$

- (b) Given a dataset $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}$, what is the negative log likelihood of the dataset according to our model? (Simplify.)

Your answer: $p(y^{(i)} | x^{(i)}, w) = \frac{1}{2} \log\left(\frac{1}{\sqrt{2\pi}}\right) \sum_i^{|\mathcal{D}|} (y^{(i)} - w x^{(i)})^2$